

ABHINAV KAUSHIK

Senior AI Engineer

New Delhi, India | +91 8700571859 | aabhinav.kaushik@gmail.com | LinkedIn

SUMMARY

Data scientist with 8+ years across the data science lifecycle in Healthcare, Telecom and AIOps, with an MSc in Machine Learning and AI. Builds forecasting, prediction and anomaly detection models and agentic AI with LangGraph, from feature engineering and validation to deployment with MLflow and Docker on Azure and GCP. Has delivered POCs and consulting for 15+ clients, explaining data clearly to technical and non-technical audiences. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group

Bengaluru, India | Nov 2025 - Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents and automating the work behind 800 Jira tickets and PRD tasks.
- Engineered a production agentic AI analytics chatbot in LangGraph for AARP medical plans, answering business questions with PySpark analysis, multi-turn reasoning and live charts, cutting manual analysis by 98%.
- Built automated evaluation and observability pipelines for Question Answering and Summary Agents with LLM evaluators, tracing and failure detection, reducing human-in-the-loop review by 43%.
- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.

Senior Machine Learning Engineer | Viavi Solutions

Chennai, India | Sep 2024 - Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients in cross-functional teams, turning client challenges into products with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions

Bengaluru, India | Nov 2021 - Sep 2024

- Deployed a production SLA breach prediction model with LightGBM, FastAPI and MLflow, monitoring performance and using GPT-3.5 for real-time root-cause analysis per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive planning for clients.
- Delivered technical POCs and client consulting for 15+ clients across 3+ AIOps and ML platforms, demonstrating product capabilities and contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum

Gurugram, India | Jun 2021 - Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems

Gurugram, India | Nov 2019 - May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft

Gurugram, India | Dec 2017 - Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree | ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

Price Elasticity and Demand Forecasting

- Built a reusable pricing and demand forecasting product estimating price elasticity per product segment with LightGBM and PyTorch, keeping feature design configurable across different client datasets.

- Validated forecasts against Holt-Winters baselines on rolling backtests, explained price drivers to business users, and deployed on Azure with MLflow tracking and Git-based CI.
- Technologies:** Python, PyTorch, LightGBM, SQL, Azure, MLflow, Git, Docker

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection | Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer | Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 - Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 - Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract
Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning
Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis
Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix
MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel
Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase
Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB
Search: Apache Solr, Lucene Indexing, Information Retrieval
BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies