

ABHINAV KAUSHIK

Data Scientist

Bengaluru, India | +91 87005 71859 | aabhinav.kaushik@gmail.com | LinkedIn

SUMMARY

Data Scientist with 8+ years turning product and business questions into statistical and machine-learning models across Healthcare, Telecom and AIOps, with an MSc in Machine Learning and AI. Works in SQL and Python on feature engineering, predictive and diagnostic modelling, and validation that stays honest about leakage and uncertainty. Owns models end to end, from outcome definition to operationalization and monitoring, and partners with stakeholders so findings turn into decisions. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group

Bengaluru, India | Nov 2025 - Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents and automating the work behind 800 Jira tickets and PRD tasks.
- Engineered a production multi-agent analytics chatbot in LangGraph for medical plans, turning stakeholders' business questions into PySpark analysis, multi-turn reasoning and live charts, cutting manual analysis by 98%.
- Built a mock-generator pipeline for customer delivery, producing 36+ discipline-specific UI mockups from Jira stories and requirements using HTML, Next.js and Claude 4.6 to accelerate solution prototyping and stakeholder review.
- Built automated evaluation and observability pipelines for Question Answering and Summary Agents with LLM evaluators, tracing and failure detection, reducing human-in-the-loop review by 43%.
- Fine-tuned a cross-encoder on positive/negative pairs to re-rank AARP survey retrieval, using Databricks notebooks and Vector Search for embeddings and retrieval, improving downstream question-answering accuracy.

Senior Machine Learning Engineer | Viavi Solutions

Chennai, India | Sep 2024 - Nov 2025

- Deployed a production SLA breach prediction model with LightGBM, FastAPI and MLflow, operationalized for recurring use with GPT-3.5 giving root-cause analysis and remediation per breach, reducing downtime by 67%.
- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built and deployed an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients across NAM, EMEA and LATAM from client requirements, with 7 moving into production or paid client work.
- Built CI/CD pipelines for model deployment with Jenkins, Bitbucket and Docker, adding unit and integration tests and automated checks so model releases stayed repeatable and reviewed before reaching production.

Data Scientist | Viavi Solutions

Bengaluru, India | Nov 2021 - Sep 2024

- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices with Isolation Forest and Holt-Winters, deriving robust signals that drove earlier detection and proactive planning.
- Contributed to network-device inventory demand forecasting using predictive models to improve planning and reduce revenue-at-risk; targeted a conservative 10% reduction in excess/at-risk inventory and higher ARPU.
- Trained and deployed a production time-series forecasting framework combining LSTM and XGBoost to model call quality, SMS delivery rates and network efficiency metrics for the Delta cellular network.
- Engineered a unified time-series data pipeline in SQL and Python to ingest call logs, SMS delivery statistics and signal metrics across multiple time intervals, joined by user ID and network location for downstream forecasting.

Data Analyst | Optum

Gurugram, India | Jun 2021 - Nov 2021

- Optimized a scikit-learn Random Forest for member risk stratification with 12+ engineered features, GridSearchCV and 5-fold cross-validation, lifting AUC from 0.74 to 0.81 and recall from 0.62 to 0.73.
- Optimized hospital staff scheduling in Python as a mixed-integer linear program with shift-coverage constraints, and modelled patient wait times with queueing statistics to balance staffing levels against demand.
- Analyzed member plan utilization and renewal behaviour across plan segments in SQL on BigQuery, testing assumptions about the drivers and building reusable segmentation queries that fed downstream risk modelling.

Data Engineer | Nippon Data Systems

Gurugram, India | Nov 2019 - May 2021

- Engineered SQL, PySpark and Airflow pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft

Gurugram, India | Dec 2017 - Nov 2019

- Built customer dashboards for financial and product insights using QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

- Developed an XGBoost propensity model scoring the likelihood of a personal-loan application within 7 days of contact, giving commercial teams a clear signal for targeted outreach.
- Wrote complex SAP HANA SQL with CTEs, joins, subqueries, window functions and aggregations to extract and validate financial modelling data.
- Built reusable SQL data-preparation workflows across financial and product datasets, refining queries and filters to surface relevant signals for financial teams.

PROJECTS

Think Tree | ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

Product Usage Churn Scoring Engine

- Built a churn and renewal risk scoring engine for a desktop software product, engineering behavioural usage features from licence and telemetry data with LightGBM and logistic regression.
- Validated it with time-based splits and leakage checks, calibrated the scores with uncertainty bands, and operationalized recurring scoring in Airflow with MLflow tracking and drift monitoring for commercial teams.
- **Technologies:** SQL, Python, pandas, scikit-learn, LightGBM, Airflow, MLflow, Power BI

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection | Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer | Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 - Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 - Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning, A/B Testing

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, CI/CD, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Databricks Vector Search, BigQuery, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native
German: A1, Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies