

ABHINAV KAUSHIK

Senior Generative AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Generative AI Engineer | [at] Group

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Generative AI Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Data Scientist | [at] Group

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Engineer - Data & AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Engineer - Data & AI | A.P. Moller - Maersk

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Engineer - Data & AI role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Data Scientist | A11

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Data Scientist role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist - Search

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist - Search | ABOUT YOU

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist - Search role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer for Generative AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer for Generative AI | Accenture DACH

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer for Generative AI role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI & Data Science Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI & Data Science Consultant | Accenture DACH

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI & Data Science Consultant role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Desarrollador senior IA Generativa - LLM y Machine Learning

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Desarrollador senior IA Generativa - LLM y Machine Learning | Accenture España

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Desarrollador senior IA Generativa - LLM y Machine Learning role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist/Data Engineer Analyst/Consultant for Supply Chain

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist/Data Engineer Analyst/Consultant for Supply Chain | Accenture España

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist/Data Engineer Analyst/Consultant for Supply Chain role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Generative AI Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Generative AI Consultant | Accenture Italia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Generative AI Consultant role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI/ML Computational Scientist Manager

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI/ML Computational Scientist Manager | Accenture Italia

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI/ML Computational Scientist Manager role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Vil du være med på å utvikle AI-løsninger som brukes av tusenvis av mennesker?

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Vil du være med på å utvikle AI-løsninger som brukes av tusenvis av mennesker? | Accenture Nordics

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Vil du være med på å utvikle AI-løsninger som brukes av tusenvis av mennesker? role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI-utvikler søkes – bli med på å forme morgendagens teknologi

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI-utvikler søkes – bli med på å forme morgendagens teknologi | Accenture Nordics

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI-utvikler søkes – bli med på å forme morgendagens teknologi role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Accenture Poland

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | ACT-ON DATA

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | ADMIRAL Technologies

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | agap2 Spain

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Engineer | Agile Robots SE

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Ahold Delhaize

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Developer – B2B

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Developer – B2B | ALGOTEQUE Innovation Hub

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Developer – B2B role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | AllCloud

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - European Agency

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - European Agency | Altia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - European Agency role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Applied Scientist, Amazon Compliance and Safety Services

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Applied Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Applied Scientist, Amazon Compliance and Safety Services | Amazon

I am a Senior Applied Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Applied Scientist, Amazon Compliance and Safety Services role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Applied Scientist, Amazon Compliance and Safety Services

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Applied Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Applied Scientist, Amazon Compliance and Safety Services | Amazon Science

I am a Senior Applied Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Applied Scientist, Amazon Compliance and Safety Services role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer / KI-Spezialist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer / KI-Spezialist | ANKÖ - Auftragnehmerkataster Österreich

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer / KI-Spezialist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer | Anton Paar

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Full-Stack Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Full-Stack Machine Learning Engineer | Aptic AB

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Full-Stack Machine Learning Engineer role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI / Machine Learning Engineer – Generative AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI / Machine Learning Engineer – Generative AI | Aramisauto

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI / Machine Learning Engineer – Generative AI role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist, FinCrime

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist, FinCrime | ARQ

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist, FinCrime role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist – MMM Expert

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist – MMM Expert | Artefact

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist – MMM Expert role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - GenAI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist - GenAI | Artefact

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - GenAI role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer / Technical Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer / Technical Consultant | Artefact

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer / Technical Consultant role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Partner AI Strategy & Transformation

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Partner AI Strategy & Transformation | artefactlinkedin

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Partner AI Strategy & Transformation role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

CDI - AI Designer Motion & Génératif - Artefact 3000

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

CDI - AI Designer Motion & Génératif - Artefact 3000 | artefactlinkedin

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the CDI - AI Designer Motion & Génératif - Artefact 3000 role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Asana

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Operations Research Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Operations Research Engineer | aspaara AG

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Operations Research Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer - Praca zdalna

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer - Praca zdalna | Astek

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer - Praca zdalna role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - Warszawa, Polska K/M

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist - Warszawa, Polska K/M | Astek

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - Warszawa, Polska K/M role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Austrian Lotteries

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist & AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist & AI Engineer | Austrian Lotteries

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist & AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Data Scientist | AutoScout24

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | AutoScout24

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Solutions Architect - Data & AI Engineer, German Speaking

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Solutions Architect - Data & AI Engineer, German Speaking | Auvaria

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Solutions Architect - Data & AI Engineer, German Speaking role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Research Scientist - Frontier Industrial AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Research Scientist - Frontier Industrial AI | AVEVA

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Research Scientist - Frontier Industrial AI role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Global Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Global Data Scientist | Avolta

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Global Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist | Avvale

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Staff/Principal Software Engineer – AI Applications

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Staff/Principal Software Engineer – AI Applications | Axelera AI

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Staff/Principal Software Engineer – AI Applications role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer till Axfood IT, Stockholm

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer till Axfood IT, Stockholm | Axfood

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer till Axfood IT, Stockholm role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Axpe Consulting

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Ayesa Digital

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | B3 Consulting Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist - Pricing

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist - Pricing | Back Market

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist - Pricing role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Ekspert / Ekspertka ds. Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Ekspert / Ekspertka ds. Data Science | Bank Millennium

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Ekspert / Ekspertka ds. Data Science role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | BAUMLINK

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Bayer

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer-1

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer-1 | Bayer

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer-1 role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Germany & Austria - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Germany & Austria - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Germany & Austria - BCG X role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Engineer, Spain - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Engineer, Spain - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Engineer, Spain - BCG X role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Germany & Austria - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Germany & Austria - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Germany & Austria - BCG X role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Italy - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Italy - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Italy - BCG X role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Germany & Austria - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Germany & Austria - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Germany & Austria - BCG X role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, Operations & Optimization, Germany & Austria - BCG X role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Scientist, France - BCG X

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Scientist, France - BCG X | BCG X

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Scientist, France - BCG X role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Vil du utvikle data- og AI-løsninger som skaper verdi?

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Vil du utvikle data- og AI-løsninger som skaper verdi? | BDO Norge

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Vil du utvikle data- og AI-løsninger som skaper verdi? role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data & AI Lead

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data & AI Lead | BearingPoint

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data & AI Lead role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist – Senior Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist – Senior Consultant | BearingPoint

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist – Senior Consultant role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist med fokus på Industry 4.0

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist med fokus på Industry 4.0 | BearingPoint

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist med fokus på Industry 4.0 role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | BearingPoint

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Job / GenAI - Data Scientist / Brussel

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Job / GenAI - Data Scientist / Brussel | Belfius

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Job / GenAI - Data Scientist / Brussel role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Developer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Developer | BETWEEN Group

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Developer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist – BIP AI LAB

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist – BIP AI LAB | BIP

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist – BIP AI LAB role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | BizAway

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Associate - Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Associate - Data Science | Black Swans Exist

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Associate - Data Science role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | BNP Paribas Fortis

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Staff Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Staff Data Scientist | Bolt

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Staff Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Booz Allen Hamilton

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI/ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI/ML Engineer | bridge351

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI/ML Engineer role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer Agentic AI Systems

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer Agentic AI Systems | Brillio

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer Agentic AI Systems role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Applied Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Applied Scientist | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Applied Scientist | bsport

I am a Senior Applied Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Applied Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist | BULLIT MANAGEMENT SERVICES LIMITED

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Bunge

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant Data Management & Machine Learning

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant Data Management & Machine Learning | BWI GmbH

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant Data Management & Machine Learning role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant Data Management & Machine Learning

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant Data Management & Machine Learning | BWI GmbH

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant Data Management & Machine Learning role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / Senior Data Scientist | C3 AI

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / Senior Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - Defence

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - Defence | C3 AI

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - Defence role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Consultant | Capgemini

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Consultant role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI & Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI & Data Scientist | Capgemini

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI & Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant Agentic AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant Agentic AI | Capgemini

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant Agentic AI role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Capgemini Engineering

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | Capitole

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer / AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859
LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer / AI Engineer | Capitole

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer / AI Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

ML/AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

ML/AI Engineer | CGI

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the ML/AI Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior ML Engineer, AI Security

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior ML Engineer, AI Security | Check Point Software

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior ML Engineer, AI Security role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist I

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist I | Checkout.com

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist I role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Ciklum

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Expert AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Expert AI Engineer | Ciklum

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Expert AI Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI/ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI/ML Engineer | Circle K

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI/ML Engineer role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist and Model Developer fixed term contract till end of Dec 2027, Vice President

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist and Model Developer fixed term contract till end of Dec 2027, Vice President | Citi

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist and Model Developer fixed term contract till end of Dec 2027, Vice President role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Clariness

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Clarivate

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Cofidis España

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Cognite

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Agentic AI Architect

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Agentic AI Architect | Cognizant

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Agentic AI Architect role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer - Center of Excellence

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer - Center of Excellence | Columbus

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer - Center of Excellence role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | Comply365

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI engineer | Cream Consulting

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI engineer role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Sr. Data Scientist - Gen AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Sr. Data Scientist - Gen AI | CrowdStrike

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Sr. Data Scientist - Gen AI role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

R&D Engineer Industrial AI 100%

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

R&D Engineer Industrial AI 100% | CSEM

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the R&D Engineer Industrial AI 100% role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Curamando

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Daintta

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Data Wizards

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Forward Deployed Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Forward Deployed Engineer | Databricks

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Forward Deployed Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Dataiku

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | Datatonic

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Datisation

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist – Machine Deep learning

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist – Machine Deep learning | Davidson consulting

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist – Machine Deep learning role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Lead

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Lead | DCC Energi A/S

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Lead role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultor/a de AI especializado/a en Generative AI & Azure

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultor/a de AI especializado/a en Generative AI & Azure | Deloitte

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultor/a de AI especializado/a en Generative AI & Azure role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Experienced - AI Engineer - NextHub Bari

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Experienced - AI Engineer - NextHub Bari | Deloitte

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Experienced - AI Engineer - NextHub Bari role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist AI, DHL Data & AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist AI, DHL Data & AI | DHL

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist AI, DHL Data & AI role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Discover International

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Member of Technical Staff - AI/ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Member of Technical Staff - AI/ML Engineer | Domera Labs

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Member of Technical Staff - AI/ML Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | DT

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI & Data Science Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI & Data Science Consultant | DXC Technology

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI & Data Science Consultant role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Architect

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

AI Architect | E-Solutions

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Architect role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Staff AI/ML Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Staff AI/ML Scientist | EcoVadis

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Staff AI/ML Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Staff AI/ML Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Staff AI/ML Scientist | EcoVadis

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Staff AI/ML Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | eGlue

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist I

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist I | Elsevier

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist I role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Elsevier

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist Senior

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist Senior | emagine

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist Senior role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Science/AI Specialist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Science/AI Specialist | Endava

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Science/AI Specialist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Ingénieur Data Science confirmé

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Ingénieur Data Science confirmé | Epsilon France

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Ingénieur Data Science confirmé role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer | Eraneos

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed Data & AI Engineer 80-100%

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed Data & AI Engineer 80-100% | Eraneos

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed Data & AI Engineer 80-100% role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Engineer / f/m/d

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Engineer / f/m/d | ERGO

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Engineer / f/m/d role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Entry Level: Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Entry Level: Data Scientist | Ericsson

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Entry Level: Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Esprinet Italia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / GenAI / CDI / H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / GenAI / CDI / H/F | Eulidia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / GenAI / CDI / H/F role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer – LLM Systems

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer – LLM Systems | Evollabs Tech

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer – LLM Systems role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Expert – Generative AI & Agentic Workflows

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Expert – Generative AI & Agentic Workflows | Experts Inside AG

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Expert – Generative AI & Agentic Workflows role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Artificial Intelligence Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Artificial Intelligence Engineer | Extia

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Artificial Intelligence Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Engineer - AI & Data

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Engineer - AI & Data | EY

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Engineer - AI & Data role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer | Fabasoft

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Data Scientist | Faculty

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Faculty

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - Deep Learning Forecasting

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production ML, anomaly detection and time-series forecasting systems across Healthcare, Telecom and AIOps. Works with LSTM, XGBoost, Holt-Winters, Python, SQL and PySpark, alongside production APIs and MLflow. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Developed time-series forecasting of base-station call handover and SMS handover volumes using LSTM and XGBoost.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters.
- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, reducing downtime by 67%.
- Built a forecasting feature for the AIOps monitoring product, forecasting univariate KPIs across the network.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - Deep Learning Forecasting | Fetcherr

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - Deep Learning Forecasting role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

F1 Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

F1 Data Scientist | FIA – Fédération Internationale de l'Automobile

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the F1 Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior / Lead Data Scientist – Applied Statistics & Machine Learning – €100,000 – Hybrid – Frankfurt

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior / Lead Data Scientist – Applied Statistics & Machine Learning – €100,000 – Hybrid – Frankfurt | Findr

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior / Lead Data Scientist – Applied Statistics & Machine Learning – €100,000 – Hybrid – Frankfurt role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior / Lead Data Scientist – Applied Statistics & Machine Learning – €100,000 – Hybrid

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior / Lead Data Scientist – Applied Statistics & Machine Learning – €100,000 – Hybrid | Findr

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior / Lead Data Scientist – Applied Statistics & Machine Learning – €100,000 – Hybrid role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | Finom

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist Senior

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist Senior | Fives

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist Senior role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Flower

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Engineer og AI-utvikler

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Engineer og AI-utvikler | Forte

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Engineer og AI-utvikler role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Fronius Österreich

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist til vores Machine Learning team

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist til vores Machine Learning team | Fødevarestyrelsen

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist til vores Machine Learning team role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Machine Learning Scientist, Frontier Research, AI for Drug Discovery

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Machine Learning Scientist, Frontier Research, AI for Drug Discovery | Genentech

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Machine Learning Scientist, Frontier Research, AI for Drug Discovery role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | GENOS

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Prisanalytiker/Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Prisanalytiker/Data Scientist | Gjensidige

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Prisanalytiker/Data Scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Developer & Data Scientist IRC303082

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Developer & Data Scientist IRC303082 | GlobalLogic

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Developer & Data Scientist IRC303082 role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Edge AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Edge AI Engineer | Grab

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Edge AI Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist - Real-Time Esports Predictions

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist - Real-Time Esports Predictions | GRID esports

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist - Real-Time Esports Predictions role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | Harmattan AI

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - AI | HCLTech

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - AI role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Solutions Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Solutions Engineer | hello again

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Solutions Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist, Growth Alliance

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist, Growth Alliance | HelloFresh

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist, Growth Alliance role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist, Demand Forecasting, Food Systems

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production ML, anomaly detection and time-series forecasting systems across Healthcare, Telecom and AIOps. Works with LSTM, XGBoost, Holt-Winters, Python, SQL and PySpark, alongside production APIs and MLflow. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Developed time-series forecasting of base-station call handover and SMS handover volumes using LSTM and XGBoost.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters.
- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, reducing downtime by 67%.
- Built a forecasting feature for the AIOps monitoring product, forecasting univariate KPIs across the network.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist, Demand Forecasting, Food Systems | HelloFresh

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist, Demand Forecasting, Food Systems role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer.

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer. | Helvetia Insurance Austria

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer. role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Artificial Intelligence Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Artificial Intelligence Engineer | Hexagon AB

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Artificial Intelligence Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / KYC @ING Hubs Romania

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / KYC @ING Hubs Romania | Hipo.ro

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / KYC @ING Hubs Romania role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data & AI Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data & AI Consultant | HiQ

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data & AI Consultant role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer - Computer Vision

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer - Computer Vision | Hire Feed

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer - Computer Vision role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Process Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Process Data Scientist | Holcim

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Process Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Process Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Process Data Scientist | Holcim Supplementary Pension Fund

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Process Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant Data / AI X/ F/H

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant Data / AI X/ F/H | HOUSE OF ABY

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant Data / AI X/ F/H role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI /ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI /ML Engineer | HP

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI /ML Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Expert Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Expert Data Science | HSE - Home Shopping Europe GmbH

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Expert Data Science role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Scientist | HUK-COBURG

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Scientist role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Artificial Intelligence Specialist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Artificial Intelligence Specialist | Humans.tech

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Artificial Intelligence Specialist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Huxley

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Amsterdam. My fiancée lives in Berlin, and I am building my career in Amsterdam while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Applied Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Applied Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Applied Scientist | i3

I am a Senior Applied Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Applied Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | If Insurance

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist | If Insurance

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | If Insurance

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Software Engineer, AI Agents

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Software Engineer, AI Agents | IFS

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Software Engineer, AI Agents role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | In4Matic

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Staff Specialist Artificial Intelligence)

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Staff Specialist Artificial Intelligence) | Infineon Technologies

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Staff Specialist Artificial Intelligence) role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Architect

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Architect | Infosys

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Architect role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / KYC @ING Hubs Romania

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / KYC @ING Hubs Romania | ING Hubs Romania

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / KYC @ING Hubs Romania role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist & ML Engineer - Predictive Analytics & Agentic AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist & ML Engineer - Predictive Analytics & Agentic AI | Intersnack IT KG

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist & ML Engineer - Predictive Analytics & Agentic AI role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Computer Vision & Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Computer Vision & Machine Learning Engineer | Intuitive

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Computer Vision & Machine Learning Engineer role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Solutions Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Solutions Engineer | IQVIA

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Solutions Engineer role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer / Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer / Data Scientist | IRiS Grupo Tecnológico

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer / Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Developer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Developer | IT Factory

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Developer role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed AI Engineer - GenAI / Agentic AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed AI Engineer - GenAI / Agentic AI | ITERIS

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed AI Engineer - GenAI / Agentic AI role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist | IXOPAY

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist H/F | Jems Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist H/F role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Software Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Software Engineer | Jet HR

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Software Engineer role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Jet HR

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI/ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI/ML Engineer | Jobloom

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI/ML Engineer role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Ingénieur IA/ML

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Ingénieur IA/ML | Jobloom

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Ingénieur IA/ML role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | Joppy

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Ops Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Ops Engineer | Journal

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Ops Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Just Eat Takeaway.com

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Amsterdam. My fiancée lives in Berlin, and I am building my career in Amsterdam while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist GenAI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist GenAI | Kaino

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist GenAI role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

KI-Engineer / KI-Plattform-Spezialist/in

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

KI-Engineer / KI-Plattform-Spezialist/in | Kanton Luzern

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the KI-Engineer / KI-Plattform-Spezialist/in role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Ingénieur Data scientist IA H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Ingénieur Data scientist IA H/F | Kicklox

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Ingénieur Data scientist IA H/F role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | KICKS

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist - Fraud Model Validation

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist - Fraud Model Validation | Klarna

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist - Fraud Model Validation role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer/Lead AI Developer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer/Lead AI Developer | Kompetenzzentrum DOK GmbH

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer/Lead AI Developer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Expert:in

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Expert:in | KPMG Austria

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Expert:in role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Solutions & Adoption Specialist*

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Solutions & Adoption Specialist* | KTM AG

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Solutions & Adoption Specialist* role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist | Körber

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Körber Pharma

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Generative AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Generative AI Engineer | La Fosse

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Generative AI Engineer role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | La Fosse

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Vision System Data Science Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Vision System Data Science Engineer | Lam Research

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Vision System Data Science Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist IA – Développeur Machine Learning H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist IA – Développeur Machine Learning H/F | LCL

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist IA – Développeur Machine Learning H/F role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Engineer | Lighthouse

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Data Scientist | Lutech

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Data Science Engineer / Bucharest

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Data Science Engineer / Bucharest | Luxoft

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Data Science Engineer / Bucharest role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Data Science Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Data Science Engineer | Luxoft

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Data Science Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | LVMH

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Macadam

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

GenAI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

GenAI Engineer | Machine Learning Reply

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the GenAI Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Machine Learning Reply

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Agentic AI Software Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Agentic AI Software Engineer | Machine Learning Reply

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Agentic AI Software Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Consultant | Machine Learning Reply

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Consultant role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Machine Learning Reply GmbH

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Consultant | Machine Learning Reply GmbH

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Consultant role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Machine Learning Reply GmbH

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Consultant | Machine Learning Reply GmbH

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Consultant role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Advanced Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Advanced Machine Learning Engineer | Magna International

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Advanced Machine Learning Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI/ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI/ML Engineer | Magnet Forensics

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI/ML Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

CONSULTOR DATA SCIENCE BARCELONA | Management Solutions

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the CONSULTOR DATA SCIENCE BARCELONA role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | MANGO

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Science Analyst - OptimusAI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Science Analyst - OptimusAI | McKinsey & Company

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Science Analyst - OptimusAI role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - McKinsey Transformation

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - McKinsey Transformation | McKinsey & Company

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - McKinsey Transformation role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Machine Learning Engineer | MedHelp Care

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist | Mindera

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Innovation Engineer Automatisiertes Fahren

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Innovation Engineer Automatisiertes Fahren | Mitsui Chemicals Group

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Innovation Engineer Automatisiertes Fahren role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer Automatisiertes Fahren

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer Automatisiertes Fahren | Mitsui Chemicals Group

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer Automatisiertes Fahren role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Specialist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Specialist | MjAMjAM Petfood GmbH

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Specialist role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI/ML Engineer Autonomous Systems

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI/ML Engineer Autonomous Systems | MOTOR Ai

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI/ML Engineer Autonomous Systems role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Mutualité chrétienne (MC)

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Voice AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Voice AI Engineer | N-iX

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Voice AI Engineer role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Consultant | NAVAX

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Consultant role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Scientist | NearTech Search

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Staff / Principal Applied AI Researcher

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Staff / Principal Applied AI Researcher | Nebius

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Staff / Principal Applied AI Researcher role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Applied ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Applied ML Engineer | Nebius

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Applied ML Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Pod Lead

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Pod Lead | Neko Health

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Pod Lead role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Pod Lead

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Pod Lead | Neko Health

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Pod Lead role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Netconomy

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Nexer Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Nexer Insight

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead IA GENERATIVE specialist - F/H

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead IA GENERATIVE specialist - F/H | Niji

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead IA GENERATIVE specialist - F/H role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | NOBA Bank Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior & Expert Data Scientists - Build the Future of AI in Operations

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior & Expert Data Scientists - Build the Future of AI in Operations | Nordea

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior & Expert Data Scientists - Build the Future of AI in Operations role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior & Expert Data Scientists - Build the Future of AI in Operations

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior & Expert Data Scientists - Build the Future of AI in Operations | Nordea

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior & Expert Data Scientists - Build the Future of AI in Operations role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist, AI platform

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist, AI platform | Nordnet Bank AB

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist, AI platform role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data scientist | Norsk helsenett SF

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior KI utvikler og data scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior KI utvikler og data scientist | Norsk helsenett SF

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior KI utvikler og data scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

PhD Candidate in AI-assisted medical imaging with MRI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

PhD Candidate in AI-assisted medical imaging with MRI | Norwegian University of Science and Technology (NTNU)

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the PhD Candidate in AI-assisted medical imaging with MRI role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Software Engineer - AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Software Engineer - AI | NOV

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Software Engineer - AI role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist IA Générative H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist IA Générative H/F | NSI

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist IA Générative H/F role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI & GenAI Architect in Innsbruck/Graz/Wien

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior GenAI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI & GenAI Architect in Innsbruck/Graz/Wien | NTT DATA Europe & Latam

I am a Senior GenAI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI & GenAI Architect in Innsbruck/Graz/Wien role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist Senior - GenAI Solutions

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist Senior - GenAI Solutions | NTT DATA Europe & Latam

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist Senior - GenAI Solutions role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | NTT DATA Europe & Latam

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Research Scientist, AI RL and LLMs

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Applied Scientist | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Research Scientist, AI RL and LLMs | NVIDIA

I am a Senior Applied Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Research Scientist, AI RL and LLMs role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

TechLead ML/CV Developer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

TechLead ML/CV Developer | ODS Serbia

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the TechLead ML/CV Developer role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist / ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist / ML Engineer | ODS Serbia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist / ML Engineer role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | OPSWAT

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Medior/Senior AI Engineer - Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Medior/Senior AI Engineer - Consultant | Orange Business

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Medior/Senior AI Engineer - Consultant role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Ekspert / Ekspertka ds. AI w Dziale Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Ekspert / Ekspertka ds. AI w Dziale Data Science | ORLEN S.A.

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Ekspert / Ekspertka ds. AI w Dziale Data Science role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Ekspert / Ekspertka ds. ML/AI w Dziale Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Ekspert / Ekspertka ds. ML/AI w Dziale Data Science | ORLEN S.A.

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Ekspert / Ekspertka ds. ML/AI w Dziale Data Science role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Design, develop & deploy ML/LLM models, data pipelines & agentic AI solutions for int/ext projects

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Design, develop & deploy ML/LLM models, data pipelines & agentic AI solutions for int/ext projects | OUTSCALE

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Design, develop & deploy ML/LLM models, data pipelines & agentic AI solutions for int/ext projects role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist Senior H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist Senior H/F | Oventi

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist Senior H/F role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist H/F

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist H/F | Oventi

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist H/F role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Oxigent Technologies

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist Applied AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist Applied AI | paretos

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist Applied AI role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Partners Group

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Peter Park

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Analyst / Data Scientist & Digital Solutions Specialist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Analyst / Data Scientist & Digital Solutions Specialist | Pfizer

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Analyst / Data Scientist & Digital Solutions Specialist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Phoenix Technologies

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Associate Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Associate Data Scientist | Predium

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Associate Data Scientist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Applied AI Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Applied AI Scientist | Preply

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Applied AI Scientist role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Software / Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Software / Machine Learning Engineer | Prima

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Software / Machine Learning Engineer role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist GenAI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist GenAI | Prima

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist GenAI role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Procter & Gamble

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - Senior Manager - Energy & Commodities

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - Senior Manager - Energy & Commodities | Publicis Sapient

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - Senior Manager - Energy & Commodities role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist Barcelona

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist Barcelona | PwC España

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist Barcelona role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning & Artificial Intelligence Engineer - Milano

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning & Artificial Intelligence Engineer - Milano | PwC Italy

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning & Artificial Intelligence Engineer - Milano role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Consultant - QuantumBlack, AI by McKinsey

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Consultant - QuantumBlack, AI by McKinsey | QuantumBlack, AI by McKinsey

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Consultant - QuantumBlack, AI by McKinsey role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist II - QuantumBlack, AI by McKinsey

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist II - QuantumBlack, AI by McKinsey | QuantumBlack, AI by McKinsey

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist II - QuantumBlack, AI by McKinsey role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist I - QuantumBlack, AI by McKinsey

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist I - QuantumBlack, AI by McKinsey | QuantumBlack, AI by McKinsey

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist I - QuantumBlack, AI by McKinsey role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Consultant - QuantumBlack, AI by McKinsey

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Consultant - QuantumBlack, AI by McKinsey | QuantumBlack, AI by McKinsey

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Consultant - QuantumBlack, AI by McKinsey role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant AI Engineer & Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant AI Engineer & Data Scientist | Quickscale AI

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant AI Engineer & Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant Senior AI Engineer & Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant Senior AI Engineer & Data Scientist | Quickscale AI

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant Senior AI Engineer & Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist | Raiffeisen Tech Poland

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - Get Ready

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - Get Ready | Reale Mutua Assicurazioni

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - Get Ready role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineering Expert

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineering Expert | Reckitt

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineering Expert role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist – TV-Insight

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist – TV-Insight | Red Bull

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist – TV-Insight role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior/Lead/Staff Applied Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior/Lead/Staff Applied Scientist | RELEX Solutions

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior/Lead/Staff Applied Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist / Data Analyst appliqué à la Conduite Autonome

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist / Data Analyst appliqué à la Conduite Autonome | Renault Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist / Data Analyst appliqué à la Conduite Autonome role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | RetInSight

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Scientist | RetInSight

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Scientist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Revolut

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Generalist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Generalist | Rex.zone

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Generalist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Scientist | Roche

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Computational & Applied AI - Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Computational & Applied AI - Scientist | Roche

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Computational & Applied AI - Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Roche

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Deep Learning Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Deep Learning Scientist | Safran.AI (ex Preligens)

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Deep Learning Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Deep Learning Scientist - IMINT

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Deep Learning Scientist - IMINT | Safran.AI (ex Preligens)

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Deep Learning Scientist - IMINT role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed Engineer | Salesforce

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed Engineer role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Forward Deployed Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Forward Deployed Engineer | Salesforce

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Forward Deployed Engineer role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Scrambly

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Scrambly

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Scrambly

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Senior Data Scientist | Scrambly

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | SDG Group Italy

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Searchability NS&D

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist | Sedona Digital

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | SEGULA Technologies

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | sennder

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist: France

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist: France | Shift Technology

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist: France role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Consultant | Sia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Consultant role in Amsterdam. My fiancée lives in Berlin, and I am building my career in Amsterdam while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant in Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant in Data Science | Sia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant in Data Science role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist & AI Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist & AI Consultant | Sia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist & AI Consultant role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior - Data Scientist & AI Consultant

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior - Data Scientist & AI Consultant | Sia

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior - Data Scientist & AI Consultant role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer | Siemens

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant – Agentic & Artificial Intelligence Siemens Advanta

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant – Agentic & Artificial Intelligence Siemens Advanta | Siemens

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant – Agentic & Artificial Intelligence Siemens Advanta role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Consultant in Agentic & Artificial Intelligence

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Consultant in Agentic & Artificial Intelligence | Siemens

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Consultant in Agentic & Artificial Intelligence role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist – Python & Cloud Data Platforms

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist – Python & Cloud Data Platforms | Sii Poland

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist – Python & Cloud Data Platforms role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist – AdTech & Healthcare

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist – AdTech & Healthcare | Sii Poland

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist – AdTech & Healthcare role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Consultant Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Consultant Data Science | Simon-Kucher

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Consultant Data Science role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | sizeless

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / ML Engineer | Skatteverket

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / ML Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | SLB

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | Sogeti

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Solirius Reply

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Product Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Product Engineer | Stott and May

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Product Engineer role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer - Applied ML & Research

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer - Applied ML & Research | Super

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer - Applied ML & Research role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Strategist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Strategist | supermodular.ai

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Strategist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Strategist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Strategist | supermodular.ai

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Strategist role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - On-Site Presence

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - On-Site Presence | Swiss Timing Ltd

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - On-Site Presence role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | Syntrace

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Solution Developer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Solution Developer | T-Systems Austria

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Solution Developer role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | TaDa

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Talan

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer - Data & Analytics

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer - Data & Analytics | Tango

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer - Data & Analytics role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / AI Specialist Categoria protetta

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / AI Specialist Categoria protetta | TAS

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / AI Specialist Categoria protetta role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - IA & Data Science Team

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - IA & Data Science Team | Teads

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - IA & Data Science Team role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Scientist | Teamis

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Tech Lead AI Services at Telenor AI Factory

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Tech Lead AI Services at Telenor AI Factory | Telenor

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Tech Lead AI Services at Telenor AI Factory role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Tenth Revolution Group

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Data Scientist | Tenth Revolution Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science & AI Specialist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science & AI Specialist | The Net Value

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science & AI Specialist role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer for The Tech Collective

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer for The Tech Collective | The Tech Collective

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer for The Tech Collective role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Applied Scientist, Document Understanding

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Applied Scientist, Document Understanding | Thomson Reuters

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Applied Scientist, Document Understanding role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Artificial Intelligence Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Artificial Intelligence Engineer | ThoughtLabs Belgium

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Artificial Intelligence Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Tibber

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - Tieto Tech Consulting

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - Tieto Tech Consulting | Tieto

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - Tieto Tech Consulting role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Titan OS

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Artificial Intelligence Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Artificial Intelligence Engineer | TMS Technology

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Artificial Intelligence Engineer role in Norway. My fiancée lives in Berlin, and I am building my career in Norway while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | Trackunit

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist – ML & Predictive Analytics *

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist – ML & Predictive Analytics * | Tri.Merge GmbH

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist – ML & Predictive Analytics * role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Research Engineer / Enterprise AI / Hybrid / TWE46943

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Research Engineer / Enterprise AI / Hybrid / TWE46943 | twentyAI

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Research Engineer / Enterprise AI / Hybrid / TWE46943 role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer - LLMs and On-Premise AI - Aarhus

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer - LLMs and On-Premise AI - Aarhus | Twoday

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer - LLMs and On-Premise AI - Aarhus role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Science Lead

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Science Lead | UCB

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Science Lead role in Brussels. My fiancée lives in Berlin, and I am building my career in Brussels while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Applied Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Applied Scientist | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Applied Scientist | UiPath

I am a Senior Applied Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Applied Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI Engineer | Undelucram.ro

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI Engineer role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Specialista/ka ds. Data Science

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: [linkedin.com/in/abhinav-kaushik10/](https://www.linkedin.com/in/abhinav-kaushik10/)

Specjalista/ka ds. Data Science | UNIQA Polska

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Specjalista/ka ds. Data Science role in Warsaw. My fiancée lives in Berlin, and I am building my career in Warsaw while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Unique AI

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in London. My fiancée lives in Berlin, and I am building my career in London while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist | URBAN LINKER

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Verti Versicherung AG

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Germany. My fiancée lives in Berlin, and I am building my career in Germany while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Associate Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Associate Data Scientist | VF Corporation

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Associate Data Scientist role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer - Data, ML & Applied AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer - Data, ML & Applied AI | Viaplay Group

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer - Data, ML & Applied AI role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist - AI & ML - Producto SaaS B2B

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist - AI & ML - Producto SaaS B2B | Vincle

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist - AI & ML - Producto SaaS B2B role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data Scientist, Revenue - Ads, Revenue

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data Scientist, Revenue - Ads, Revenue | Vinted

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data Scientist, Revenue - Ads, Revenue role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Principal Data Scientist, Generative Recommendations & AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building retrieval, ranking, recommendation and production ML systems across Healthcare, Telecom and AIOps. Works with cross-encoders, embeddings, semantic retrieval, Python, SQL and PySpark, alongside production APIs and ML evaluation. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Fine-tuned a cross-encoder on positive and negative pairs to re-rank dense retrieval over AARP survey data, improving retrieval accuracy.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.
- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, increasing conversion by 2%.
- Built SQL and PySpark data pipelines automating medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Principal Data Scientist, Generative Recommendations & AI | VistaPrint

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Principal Data Scientist, Generative Recommendations & AI role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | Voi Technology

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist | Volvo Financial Services

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist role in Sweden. My fiancée lives in Berlin, and I am building my career in Sweden while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI/ML Engineer with PhD

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI/ML Engineer with PhD | W Hunt España

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI/ML Engineer with PhD role in Barcelona. My fiancée lives in Berlin, and I am building my career in Barcelona while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Lead Data scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Lead Data scientist | W3Global

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Lead Data scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Manager - Lead Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Manager - Lead Data Scientist | Wavestone

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Manager - Lead Data Scientist role in Paris. My fiancée lives in Berlin, and I am building my career in Paris while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Scientist - Digital Pills

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Scientist - Digital Pills | Welyk

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Scientist - Digital Pills role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | Winnow

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Romania. My fiancée lives in Berlin, and I am building my career in Romania while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Machine Learning Engineer | World

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Machine Learning Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior AI ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior AI ML Engineer | Xcede

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior AI ML Engineer role in Berlin. My fiancée lives in Berlin, and I am building my career in Berlin while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Machine Learning Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Machine Learning Engineer | XpertDirect

I am a Senior Machine Learning Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Machine Learning Engineer role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist | XXXLdigital – Part of XXXL Group

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist role in Austria. My fiancée lives in Berlin, and I am building my career in Austria while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

AI Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

AI Engineer | Yves&Partners

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the AI Engineer role in Italy. My fiancée lives in Berlin, and I am building my career in Italy while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Data Scientist / Senior Data Scientist – Machine Learning & AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years delivering applied machine learning, analytics and production AI systems across Healthcare, Telecom and AIOps. Works across Python, SQL, PySpark, scikit-learn, LightGBM, LLM applications and ML evaluation, from stakeholder requirements through deployment. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior Data Scientist | United Health Group | Bengaluru, India Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Data Scientist / Senior Data Scientist – Machine Learning & AI | ZABEL

I am a Senior Data Scientist with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Data Scientist / Senior Data Scientist – Machine Learning & AI role in Munich. My fiancée lives in Berlin, and I am building my career in Munich while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior ML Engineer

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

- Network Intern, Airport Authority of India
- Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior ML Engineer | Zerolook

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior ML Engineer role in Switzerland. My fiancée lives in Berlin, and I am building my career in Switzerland while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.

ABHINAV KAUSHIK

Senior Data Analyst - Bioinformatics / AI

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

SUMMARY

Senior AI/ML Engineer with 8+ years building production AI, ML and LLM systems across Healthcare, Telecom and AIOps. Works with LangGraph, RAG, document extraction, LLM evaluation, Python, PySpark, FastAPI, Docker and MLflow across the AI engineering lifecycle. Delivers measurable outcomes with stakeholders. Alongside this, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD.

PROFESSIONAL EXPERIENCE

Senior AI Engineer | United Health Group | Bengaluru, India

Nov 2025 – Present

- Deployed a production lab-extraction agent with Claude Opus 4.7 and Amazon Textract, integrated with LIMS, handling about 1,800 lab documents.
- Engineered a production multi-agent analytics chatbot in LangGraph for AARP medical plans, using PySpark analysis and multi-turn reasoning, cutting manual analysis by 98%.
- Built automated LLM evaluation and observability pipelines with tracing and failure detection, reducing human-in-the-loop review by 43%.
- Developed a RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate.

Senior Machine Learning Engineer | Viavi Solutions | Chennai, India

Sep 2024 – Nov 2025

- Developed an internal RAG-based Project Discovery Assistant in LangGraph combining semantic retrieval with LLM reasoning, reaching a 90% match rate against existing organizational solutions.
- Built an event-correlation engine clustering alarms with Sentence-BERT and DBSCAN to trace the cascading effect of network alarm floods back to the root event, improving semantic alarm classification by 73.5%.
- Prototyped 15+ AI applications for AIOps and GeoPerformance clients from client requirements, with 7 moving into production or paid client work.

Data Scientist | Viavi Solutions | Bengaluru, India

Nov 2021 – Sep 2024

- Deployed a production SLA breach prediction system with LightGBM, FastAPI and MLflow, using GPT-3.5 for real-time root-cause analysis and remediation recommendations per breach, reducing downtime by 67%.
- Shipped anomaly detection and demand forecasting on 15+ KPIs across 500+ telecom network devices using Isolation Forest and Holt-Winters, enabling earlier detection and proactive network planning.
- Delivered technical POCs and consulting for 15+ clients across 3+ AIOps and ML platforms, contributing to an 11% increase in POC and mentoring revenue.

Data Analyst | Optum | Gurugram, India

Jun 2021 – Nov 2021

- Implemented a reranker for an existing recommendation system serving individual and child plan recommendations, optimizing renewal suggestions and increasing conversion by 2%.
- Developed a Python and Flask parser that automated extraction and transformation of raw data into structured datasets, feeding downstream analytics and reporting.

Data Engineer | Nippon Data Systems | Gurugram, India

Nov 2019 – May 2021

- Engineered SQL and PySpark data pipelines automating BAU medical-audit reporting on large healthcare datasets, cutting manual reporting effort by 95%.
- Wrote complex SQL queries with multi-table joins, CTEs and window functions, delivering ad-hoc reports and analyses for multiple stakeholder teams.
- Built real-time log observability for the COVID QRT team with Elasticsearch and Kibana, improving traceability efficiency by 80%.

Associate Consultant | Saksoft | Gurugram, India

Dec 2017 – Nov 2019

- Developed an XGBoost propensity model predicting personal-loan applications within 7 days of contact, enabling targeted outreach and customer prioritization.
- Built customer dashboards for financial and product insights with QlikSense, SQL and SAP HANA, increasing user engagement by 40%.

PROJECTS

Think Tree – ML Study Platform

- Built a non-profit AI-powered ML study platform for people with ADHD and context-switching challenges, designed to help users organize learning and maintain study context.

RFP Multi-Agent Proposal Automation Platform

- Built a multi-agent RFP platform where an LLM supervisor orchestrates intent detection, retrieval, proposal generation, document matching and market research, with RAG over historical proposals.

Financial Fraud Detection – Mar 2025

- Built a fraud detection and reasoning workflow on the JPMorgan Synthetic Dataset, pairing LightGBM scoring with LLM-generated explanations for interpretable financial-risk analysis.

Restaurant Feedback Analyzer – Aspect-Based Sentiment Analysis

- Built an NLP pipeline to extract review aspects and determine sentiment at the aspect level, enabling granular insights into customer experience.

EDUCATION

Master of Science, Machine Learning and AI Liverpool John Moores University	Jan 2021 – Jan 2023
Bachelor of Technology, Information Technology Jamia Hamdard	Aug 2013 – Aug 2017

TECHNICAL SKILLS

AI / GenAI: Generative AI, LLM, RAG, Agentic RAG, Agentic AI, LangGraph, LangChain, Claude, OpenAI, Embeddings, Cross-Encoder, Fine-Tuning, Hugging Face Transformers, NLP, Tool Orchestration, Amazon Textract

Machine Learning: PyTorch, Scikit-learn, XGBoost, LightGBM, Random Forest, Isolation Forest, LSTM, SBERT, DBSCAN, Time Series Forecasting, Demand Forecasting, Anomaly Detection, Statistical Modeling, Regression, Deep Learning

Evals / Observability: LLM Evals, LLM Evaluators, Agent Tracing, Observability, Elasticsearch, Kibana, Log Analysis

Programming / Data: Python, SQL, PySpark, Pandas, NumPy, FastAPI, Flask, Hive, Spark SQL, REST API, Pydantic, Unix

MLOps / Cloud: MLflow, Docker, Kubernetes, Airflow, DVC, Terraform, Jenkins, Bamboo, Bitbucket, GCP, Azure, Vercel

Data Platforms: Databricks, Databricks SQL, Apache Spark, Supabase

Databases: TimescaleDB, MySQL, Vector DB, pgvector, ArangoDB, Cosmos DB

Search: Apache Solr, Lucene Indexing, Information Retrieval

BI: Power BI, QlikSense, SAP HANA

LANGUAGES

English: Native

German: A1 – Pursuing

ADDITIONAL INFORMATION

-
- Network Intern, Airport Authority of India
 - Hadoop Intern, Ducat Technologies

ABHINAV KAUSHIK

Email: aabhinav.kaushik@gmail.com Phone: +91 87005 71859

LinkedIn: linkedin.com/in/abhinav-kaushik10/

Senior Data Analyst - Bioinformatics / AI | ZS

I am a Senior AI Engineer with 8+ years of experience building production AI, ML and LLM systems across healthcare, telecom and AIOps. My recent work spans retrieval, agentic applications, evaluation, analytics and production deployment.

I am applying for the Senior Data Analyst - Bioinformatics / AI role in Copenhagen. My fiancée lives in Berlin, and I am building my career in Copenhagen while planning our move together and settling there long term.

I would bring practical experience from production systems including multi-agent analytics in LangGraph, retrieval and reranking, LLM evaluation, forecasting, and Python/PySpark data workflows. I have also worked with FastAPI, Docker and MLflow to move solutions from development into operational use.

Alongside my professional work, I am building ThinkTree.AI, a non-profit AI learning platform used by NGO teachers to support students with ADHD. I would welcome a conversation about the role and the problems your team is solving.